

Predicting Concrete Compressive Strength From Its Components Quantities



Introduction

Concrete is the most important material in civil engineering and it is of five major elements in various proportions: cement, water, coarse aggregates, fine aggregates (i.e. sand), and air. Concrete admixtures are natural or manufactured chemicals or additives added during concrete mixing to enhance specific properties of the fresh or hardened concrete, such as workability, durability, or early and final strength. The concrete compressive strength is a function of age and ingredients. There have been several researches to develop the optimum mixture of Concrete.

Problem Statement

The Concrete research usually uses Concrete Compressive Strength Test in lab to find out the resultant strengths of different combination of concrete ingredients. This can be costly and time consuming. In this project, I will predict the compressive strength of concrete mixture measured in MPa using the historical records of previous testing and regression analysis methods.

Data Source

This dataset was taken from [UCI Repository](#).

Original Owner and Donor Prof. I-Cheng Yeh Department of Information Management Chung-Hua University, Hsin Chu, Taiwan 30067, R.O.C. e-mail:icyeh@chu.edu.tw TEL:886-3-5186511

Date Donated: August 3, 2007

The following table shows the names of the variables and their meanings

variable		meaning and measurement unit
cement		quantitative, kg in a m3 mixture, Input Variable
slag	Blast Furnace Slag, quantitative, kg in a m3 mixture,Input Variable	
flyash	Fly Ash, quantitative, kg in a m3 mixture,Input Variable	
water	Blast Furnace Slag, quantitative, kg in a m3 mixture,Input Variable	
superplasticizer	Superplasticizer, quantitative, kg in a m3 mixture,Input Variable	
coarseaggregate	Coarse Aggregate, quantitative, kg in a m3 mixture,Input Variable	
fineaggregate	Fine Aggregate, quantitative, kg in a m3 mixture,Input Variable	
age	Age quantitative, Day (1~365), Input Variable	
csMPa	Concrete compressive strength, quantitative, MPa, Output Variable	

Data Extraction & Preparation

In this section, I extract the data from a csv file, analyze it a little and prepare it for further analysis

```
import numpy
import pandas as pd
```

```
df = pd.read_csv('../data/Concrete_Data.csv')
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1030 entries, 0 to 1029
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   cement                1030 non-null  float64
1   slag                  1030 non-null  float64
2   flyash                1030 non-null  float64
3   water                 1030 non-null  float64
4   superplasticizer      1030 non-null  float64
5   coarseaggregate       1030 non-null  float64
6   fineaggregate         1030 non-null  float64
7   age                   1030 non-null  int64
8   csMPa                 1030 non-null  float64
dtypes: float64(8), int64(1)
memory usage: 72.5 KB
```

```
df.rename(columns={'flyash':'flya_ash',
                  'superplasticizer':'super_plasticizer',
                  'coarseaggregate':'coarse_aggregate',
                  'fineaggregate':'fine_aggregate',
                  'csMPa':'compressive_strength_csMPa'}, inplace=True)

df.head()
```

	cement	slag	flya_ash	water	super_plasticizer	coarse_aggregate	fine_aggregate	age	compressive_strength_csMPa
0	540.0	0.0	0.0	162.0	2.5	1040.0	676.0	28	79.99
1	540.0	0.0	0.0	162.0	2.5	1055.0	676.0	28	61.89
2	332.5	142.5	0.0	228.0	0.0	932.0	594.0	270	40.27
3	332.5	142.5	0.0	228.0	0.0	932.0	594.0	365	41.05
4	198.6	132.4	0.0	192.0	0.0	978.4	825.5	360	44.30

Checking for Missing Values in the Data

```
df.isnull().sum()
```

```
cement      0
slag        0
flya_ash    0
water       0
super_plasticizer  0
coarse_aggregate  0
fine_aggregate  0
age         0
compressive_strength_csMPa  0
dtype: int64
```

Checking for Duplicates Rows

```
df[df.duplicated(keep = False)]
```

	cement	slag	flya_ash	water	super_plasticizer	coarse_aggregate	fine_aggregate	age	compressive_strength_csMPa
72	425.0	106.3	0.0	153.5	16.5	852.1	887.1	3	33.40
77	425.0	106.3	0.0	153.5	16.5	852.1	887.1	3	33.40
80	425.0	106.3	0.0	153.5	16.5	852.1	887.1	3	33.40
83	362.6	189.0	0.0	164.9	11.6	944.7	755.8	3	35.30
86	362.6	189.0	0.0	164.9	11.6	944.7	755.8	3	35.30
88	362.6	189.0	0.0	164.9	11.6	944.7	755.8	3	35.30
91	362.6	189.0	0.0	164.9	11.6	944.7	755.8	3	35.30
95	425.0	106.3	0.0	153.5	16.5	852.1	887.1	7	49.20
100	425.0	106.3	0.0	153.5	16.5	852.1	887.1	7	49.20
103	425.0	106.3	0.0	153.5	16.5	852.1	887.1	7	49.20
106	362.6	189.0	0.0	164.9	11.6	944.7	755.8	7	55.90
109	362.6	189.0	0.0	164.9	11.6	944.7	755.8	7	55.90
111	362.6	189.0	0.0	164.9	11.6	944.7	755.8	7	55.90
118	425.0	106.3	0.0	153.5	16.5	852.1	887.1	28	60.29
123	425.0	106.3	0.0	153.5	16.5	852.1	887.1	28	60.29
126	425.0	106.3	0.0	153.5	16.5	852.1	887.1	28	60.29
129	362.6	189.0	0.0	164.9	11.6	944.7	755.8	28	71.30
132	362.6	189.0	0.0	164.9	11.6	944.7	755.8	28	71.30
134	362.6	189.0	0.0	164.9	11.6	944.7	755.8	28	71.30
137	362.6	189.0	0.0	164.9	11.6	944.7	755.8	28	71.30
141	425.0	106.3	0.0	153.5	16.5	852.1	887.1	56	64.30
146	425.0	106.3	0.0	153.5	16.5	852.1	887.1	56	64.30
149	425.0	106.3	0.0	153.5	16.5	852.1	887.1	56	64.30
152	362.6	189.0	0.0	164.9	11.6	944.7	755.8	56	77.30
155	362.6	189.0	0.0	164.9	11.6	944.7	755.8	56	77.30
157	362.6	189.0	0.0	164.9	11.6	944.7	755.8	56	77.30
160	362.6	189.0	0.0	164.9	11.6	944.7	755.8	56	77.30
164	425.0	106.3	0.0	153.5	16.5	852.1	887.1	91	65.20
169	425.0	106.3	0.0	153.5	16.5	852.1	887.1	91	65.20
172	425.0	106.3	0.0	153.5	16.5	852.1	887.1	91	65.20
175	362.6	189.0	0.0	164.9	11.6	944.7	755.8	91	79.30
177	362.6	189.0	0.0	164.9	11.6	944.7	755.8	91	79.30
179	362.6	189.0	0.0	164.9	11.6	944.7	755.8	91	79.30

182	362.6	189.0	0.0	164.9	11.6	944.7	755.8	91	79.30
801	252.0	0.0	0.0	185.0	0.0	1111.0	784.0	28	19.69
809	252.0	0.0	0.0	185.0	0.0	1111.0	784.0	28	19.69

Interesting!

It seems that some records are entered more than once. This may affect the results of the analysis!

I will keep the first record and remove its duplicate.

Dropping the Duplicate rows

```
df.drop_duplicates(keep='first', inplace=True)
```

Now, data is ready for further analysis

Exploratory Data Analysis

```
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats.mstats import normaltest
```

Checking The Distributions of the Variables (features and target)

```
for col in df.columns.tolist():
    print(f'{col.capitalize()}')
    print(f'{col.capitalize()} Statistics')
    print(df[col].describe())
    print('\n')

_, pvalue = normaltest(df[col].values)
if pvalue < 0.05:
    result = f'{col} not normally distributed'
else:
    result = f'{col} normally distributed'
print(f'{col.capitalize()} Normality Test Result')
print('P-value:',str(pvalue),', Decision: ',result)

plt.figure(figsize=(18,4), facecolor='lightblue')
plt.suptitle(f'{col.capitalize()} Distribution',size = 18, y = 1.05)

plt.subplot(1,3,1)
sns.histplot(data=df, x = col, bins=30)
plt.title('Histogram')


plt.subplot(1,3,2)
plt.title('Density', size = 18)
sns.kdeplot(data=df[col])

plt.subplot(1,3,3)
sns.boxplot(data=df, x = col)
plt.title('Boxplot')

plt.tight_layout()
plt.show()
```


Cement
Cement Statistics
count 1005.000000
mean 278.631343
std 104.344261
min 102.000000
25% 190.700000
50% 265.000000
75% 349.000000
max 540.000000
Name: cement, dtype: float64

Cement Normality Test Result
P-value: 9.809278312609514e-14 , Decision: cement not normally distributed




Slag
Slag Statistics
count 1005.000000
mean 72.043483
std 86.170807
min 0.000000
25% 0.000000
50% 20.000000
75% 142.500000
max 359.400000
Name: slag, dtype: float64

Slag Normality Test Result
P-value: 1.0885008864076576e-23 , Decision: slag not normally distributed




```
Flya_ash
Flya_ash Statistics
count    1005.000000
mean      55.536318
std       64.207969
min        0.000000
25%        0.000000
50%        0.000000
75%       118.300000
max       200.100000
Name: flya_ash, dtype: float64
```

```
Flya_ash Normality Test Result
P-value: 0.0 ,    Decision: flya_ash not normally distributed
```



```
Water
Water Statistics
count    1005.000000
mean     182.075323
std      21.339334
min      121.800000
25%      166.600000
50%      185.700000
75%      192.900000
max      247.000000
Name: water, dtype: float64
```


```
Water Normality Test Result
P-value: 0.5005793792135014 ,    Decision: water normally distributed
```



```
Super_plasticizer
Super_plasticizer Statistics
count      1005.000000
mean        6.033234
std         5.919967
min         0.000000
25%         0.000000
50%         6.100000
75%        10.000000
max        32.200000
Name: super_plasticizer, dtype: float64
```

Super_plasticizer Normality Test Result


P-value: 3.9839459643086454e-35 , Decision: super_plasticizer not normally distributed



```
Coarse_aggregate
Coarse_aggregate Statistics
count      1005.000000
mean      974.376816
std       77.579667
min       801.000000
25%       932.000000
50%       968.000000
75%      1031.000000
max      1145.000000
Name: coarse_aggregate, dtype: float64
```


Coarse_aggregate Normality Test Result

P-value: 3.448426277401913e-07 , Decision: coarse_aggregate not normally distributed




Fine_aggregate
Fine_aggregate Statistics
count 1005.000000
mean 772.688259
std 80.340435
min 594.000000
25% 724.300000
50% 780.000000
75% 822.200000
max 992.600000
Name: fine_aggregate, dtype: float64

Fine_aggregate Normality Test Result
P-value: 0.004291006583264156 , Decision: fine_aggregate not normally distributed




Age
Age Statistics
count 1005.000000
mean 45.856716
std 63.734692
min 1.000000
25% 7.000000
50% 28.000000
75% 56.000000
max 365.000000
Name: age, dtype: float64

Age Normality Test Result
P-value: 3.9929469938994235e-154 , Decision: age not normally distributed




```
Compressive_strength_csmpa
Compressive_strength_csmpa Statistics
count    1005.000000
mean      35.250378
std       16.284815
min        2.330000
25%       23.520000
50%       33.800000
75%       44.870000
max       82.600000
Name: compressive_strength_csMPa, dtype: float64
```

Compressive_strength_csmpa Normality Test Result
P-value: 3.154432677195138e-07 , Decision: compressive_strength_csMPa not normally distributed



Features Non-normality

Some beginners think (erroneously) that the normal distribution assumption of linear regression applies to their data. They might plot their response variable as a histogram and examine whether it differs from a normal distribution. Others assume that the explanatory variable must be normally-distributed. Neither is required. The normality assumption relates to the distributions of the residuals. This is assumed to be normally distributed, and the regression line is fitted to the data such that the mean of the residuals is zero.

[Source](#)

Removing Outliers From Target

It can be seen that the target contains some outliers. Removing those outliers can improve the scoring of the model. There are two methods for detecting and removing outliers. For normally distributed data I can use $(\text{mean} - 3\text{standard_deviation})$ and $(\text{mean} + 3\text{standard_deviation})$. For non normally distributed data, I can use the interquartile method. Since the data is not normally distributed, I will use the interquartile method

```
q1 = df.compressive_strength_csMPa.quantile(0.25)
q3 = df.compressive_strength_csMPa.quantile(0.75)
iqr = q3-q1
upper_lim = q3 + 1.5*iqr
lower_lim = q1 - 1.5*iqr

df = df[(df.compressive_strength_csMPa >= lower_lim) & (df.compressive_strength_csMPa <= upper_lim)]
```

Features Selection (based on statistical testing)

```
from scipy.stats import pearsonr
```

```
features = [col for col in df.columns if col != 'compressive_strength_csMPa']  
y = df['compressive_strength_csMPa']  
  
y.head()
```

```
1    61.89  
2    40.27  
3    41.05  
4    44.30  
5    47.03  
Name: compressive_strength_csMPa, dtype: float64
```


```
significant_features = []  
for col in features:  
    corr, p_value = pearsonr(df[col], df['compressive_strength_csMPa'])  
    print(f'{col} vs. compressive_strength_csMPa:\n')  
    print('Correlation Coefficient:', corr)  
    print('p_value:', p_value)  
  
    if p_value < 0.05:  
        significant_features.append(col)  
        if corr > 0:  
            print(f'There is a significant evidence of a positive relationship between {col} and Concrete Compressive Strength')  
        else:  
            print(f'There is a significant evidence of a negative relationship between {col} and Concrete Compressive Strength')  
    else:  
        print(f'There is no significant relationship between {col} and Concrete Compressive Strength')  
  
plt.figure(figsize = (6,4), facecolor='lightblue')  
plt.title(f'{col.capitalize()} vs. Compressive Strength of Concrete')  
sns.regplot(x=col, y='compressive_strength_csMPa', data=df)  
plt.show()
```

cement vs. compressive_strength_csMPa:

Correlation Coefficient: 0.4822980854149247

p_value: 3.260615343299635e-59


There is a significant evidence of a positive relationship between cement and Concrete Compressive Strength



Correlation Coefficient: 0.08746360483548354

p_value: 0.0057181044887169895

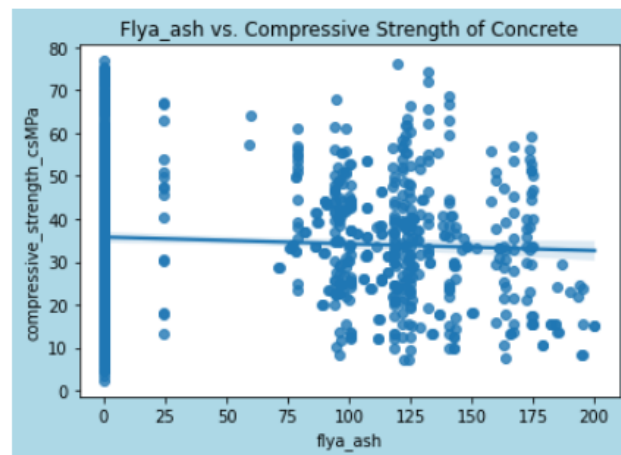
There is a significant evidence of a positive relationship between slag and Concrete Compressive Strength



Correlation Coefficient: -0.06373667178654671

p_value: 0.04421716444068777

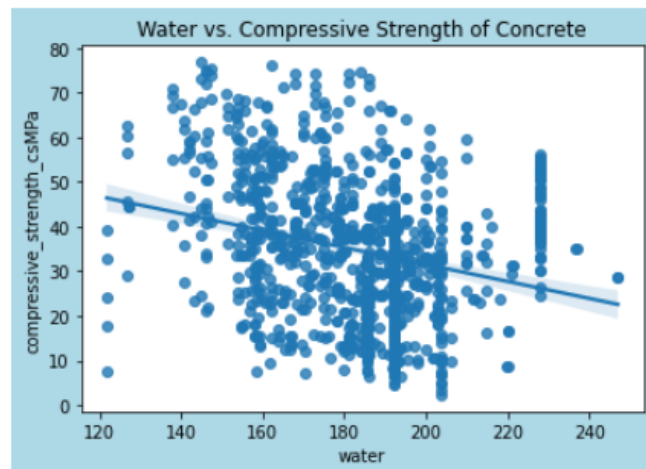
There is a significant evidence of a negative relationship between flya_ash and Concrete Compressive Strength



Correlation Coefficient: -0.2555851420974356

p_value: 2.473472671917299e-16


There is a significant evidence of a negative relationship between water and Concrete Compressive Strength



Correlation Coefficient: 0.33453422476340167

p_value: 1.704786278505969e-27

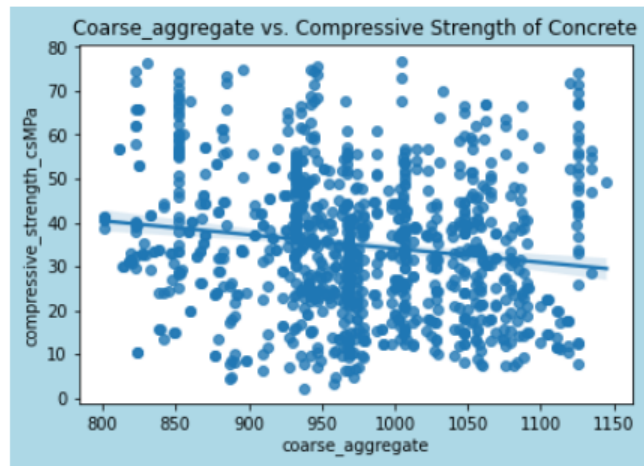
There is a significant evidence of a positive relationship between super_plasticizer and Concrete Compressive Strength



Correlation Coefficient: -0.15460377731918934

p_value: 9.339984407882778e-07

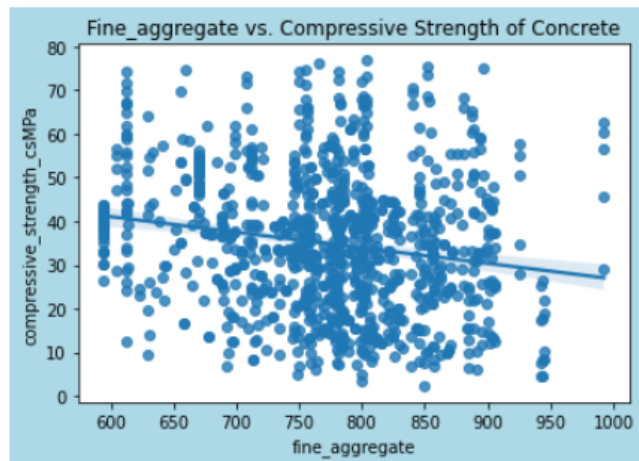
There is a significant evidence of a negative relationship between coarse_aggregate and Concrete Compressive Strength



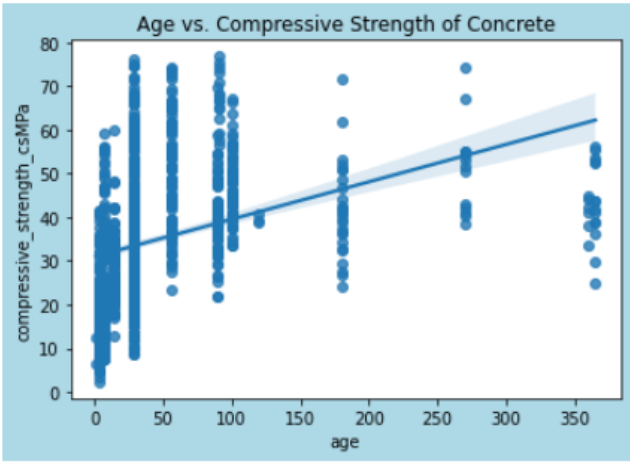
Correlation Coefficient: -0.18091803888531244

p_value: 8.771870787097434e-09

There is a significant evidence of a negative relationship between fine_aggregate and Concrete Compressive Strength



Correlation Coefficient: 0.34524351184373747
p_value: 2.752890440217549e-29
There is a significant evidence of a positive relationship between age and Concrete Compressive Strength



Dividing the Data to Features (X) and Target (y)

```
X_initial = df[significant_features]

print(X_initial.head())
```


	cement	slag	flya_ash	water	super_plasticizer	coarse_aggregate \
1	540.0	0.0	0.0	162.0	2.5	1055.0
2	332.5	142.5	0.0	228.0	0.0	932.0
3	332.5	142.5	0.0	228.0	0.0	932.0
4	198.6	132.4	0.0	192.0	0.0	978.4
5	266.0	114.0	0.0	228.0	0.0	932.0

	fine_aggregate	age
1	676.0	28
2	594.0	270
3	594.0	365
4	825.5	360
5	670.0	90

Checking for Mutlicollinearity

As a graduate of Civil Engineering (domain knowledge), I knwo there is no multicollinearity between the components of concrete mixutres as each material or additive serves different function. However, as machine learnign best practice, I should check for multicollinearity

```
corr = X_initial.corr()
plt.figure(figsize=(8,6))
sns.heatmap(corr, annot=True)
plt.xticks();
```



Looking at the values of pearson correlation, there is no significant maulticollinearity!

Feature Engineering

```
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
```

Scaling the Features

```
# It is essential for the features of a regression models to on the same scale

standarizing_transformer = StandardScaler()

X_scaled = standarizing_transformer.fit_transform(X_initial)

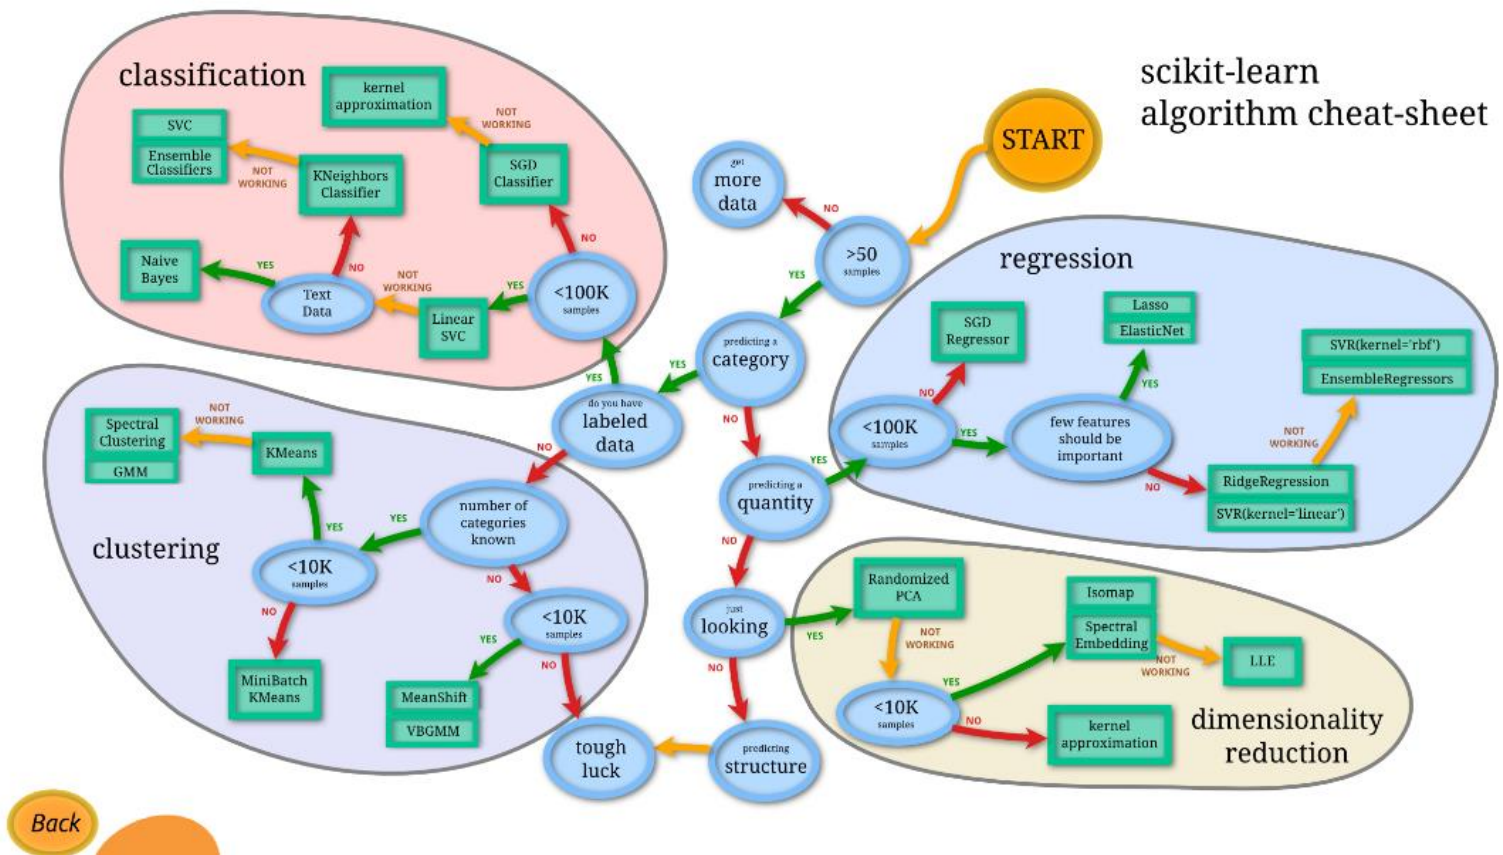
X_scaled = pd.DataFrame(X_scaled, columns = X_initial.columns)

# Although it is not necessary to normalize the targte varibale, sometimes it may lead to
# better model.
y = numpy.log1p(y)
```

[illegible]

Model Selection and Evaluation

scikit-learn
algorithm cheat-sheet



Strategy

Selecting an Estimator

- I will take the recommendation of sklearn documentation with some modification.
- I will test several baseline models initially.
- Next, I will try several ensemble regressors.
- Next, I will compare the performances of different baseline models.
- Finally, I will optimize and tune the selected model using GridSearchCV

Selecting Evaluation Scoring Metrics

I will use r2 since I don't know the permissible limit of error. The permissible limit of error differs depending on the objective of the study

```
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import KFold

from sklearn.linear_model import LinearRegression
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso

from sklearn.svm import LinearSVR

from sklearn.neighbors import KNeighborsRegressor

from sklearn.tree import DecisionTreeRegressor

from sklearn.ensemble import AdaBoostRegressor, BaggingRegressor, GradientBoostingRegressor, RandomForestRegressor

from xgboost import XGBRegressor
```

Models Building

```
l_r = LinearRegression(n_jobs=-1)

r_r = Ridge(random_state = 2)

knn_r = KNeighborsRegressor(n_jobs=-1)

dt_r = DecisionTreeRegressor(random_state = 2)

rf_r = RandomForestRegressor(random_state = 2, n_jobs= -1)

bg_r = BaggingRegressor(random_state= 2, n_jobs=-1)

xgb_r = XGBRegressor(random_state= 2, n_jobs=-1)

ada_r = AdaBoostRegressor()

gb_r = GradientBoostingRegressor()

models = [('LinearRegression',l_r),
          ('Ridge',r_r),
          ('KNeighborsRegressor',knn_r),
          ('DecisionTreeRegressor',dt_r),
          ('RandomForestRegressor',rf_r),
          ('BaggingRegressor',bg_r),
          ('AdaBoostRegressor',ada_r),
          ('GradientBoostingRegressor',gb_r),
          ('XGBRegressor',xgb_r)
        ]

results = {'model': [],
          'r2_score':[],
          'mean_absolute_error':[]
        }
```

```
for model_name, model in models:
    # X_transformed = polynomiallyTrasnformed_standarized_Labeled(X,i)

    reg_model = model

    kf = KFold(n_splits=5, shuffle= True, random_state= 2).split(X_train,y_train)
    r2_scores =cross_val_score(reg_model, X_train,y_train,cv=kf, scoring='r2')

    kf = KFold(n_splits=5, shuffle= True, random_state= 2).split(X_train,y_train)
    mae_scores =cross_val_score(reg_model, X_train,y_train,cv=kf, scoring='neg_mean_absolute_error')

    results['model'].append(model_name)
    results['r2_score'].append(r2_scores.mean())
    results['mean_absolute_error'].append(-1*mae_scores.mean())

results_df = pd.DataFrame(results)
results_df = results_df.sort_values('r2_score', ascending=False).reset_index(drop=True)
results_df
```

	model	r2_score	mean_absolute_error
0	XGBRegressor	0.914634	0.096989
1	GradientBoostingRegressor	0.904129	0.115305
2	RandomForestRegressor	0.891026	0.118985
3	BaggingRegressor	0.870394	0.129581
4	AdaBoostRegressor	0.788531	0.190277
5	DecisionTreeRegressor	0.779214	0.157721
6	KNeighborsRegressor	0.648805	0.231035
7	Ridge	0.520315	0.282283
8	LinearRegression	0.519966	0.282256

So, the model with the best r2 scoring is XGBRegressor and with no polynomial transformation

Testing the Selected Model

```
from sklearn.metrics import r2_score
```

```
model = XGBRegressor(random_state= 2, n_jobs=-1)

model.fit(X_train,y_train)

y_pred = model.predict(X_test)

print('The score of the test:')
print(r2_score(y_test,y_pred))
```

The score of the test:
0.9307441376190104

Tuning the Model

```
from sklearn.model_selection import GridSearchCV
```

```
xgb_r = XGBRegressor(random_state= 2, n_jobs=-1)

params_grid = {'eta':[0.1,0.2,0.3,0.4]}

kf = KFold(n_splits=5, shuffle= True, random_state= 2).split(X_scaled,y)

gs_xgb_r = GridSearchCV(xgb_r, params_grid, n_jobs=-1, cv = kf, scoring='r2')

gs_xgb_r.fit(X_scaled,y);

print('Best eta')
print(gs_xgb_r.best_params_)
print('\n')
print('Best score:')
print(gs_xgb_r.best_score_)
```

Best eta
{'eta': 0.2}

Best score:
0.9334140397687536

Model Interpretation

```
xgb_reg = XGBRegressor(random_state= 2, n_jobs=-1, eta = 0.2)
xgb_reg.fit(X_scaled,y)

importance_df = pd.DataFrame({'feature':X_scaled.columns.tolist(),
                             'importance':xgb_reg.feature_importances_}).\
                             sort_values('importance',ascending=False)

importance_df
```

	feature	importance
7	age	0.312057
0	cement	0.204868
4	super_plasticizer	0.182134
1	slag	0.106332
3	water	0.074790
6	fine_aggregate	0.052398
2	flya_ash	0.037724
5	coarse_aggregate	0.029698

From the importance dataframe, it can be seen that the 3 most impactfull measureable factors in the compressive strength of the concrete are age (in days), cement amount and superplasticizer amount. For each 1 day, 1 unit of cement (kg in m3) and unit of superplasticizer (kg in m3) the compressive strength of oncrete inscreases by 0.31 MPa, 0.2 MPa and 0.18 MPa respectively.

Saving the Model

```
import joblib
```

```
joblib.dump(xgb_reg, 'xgb_reg.pkl')
```

```
['xgb_reg.pkl']
```

This model can be used to estimate the compressive strength of concrete before carrying out the physical experiemnt. By experimenting with different values of (water, cement, fine aggregate...ect), a reseacher can use this model to estimate the corresponding concrete compressive strength before carrying out the actual experiment in the lab. That can save effort, time and cost.